

Data Structures and Data Frames

Preflight

1. Ensure that RStudio/R is open with the appropriate packages installed (tidyverse)
-

Learning Questions

- This lesson might sound like it's going to be a bit dry: "Why should I care about data types and structures?"
 - **IF YOU UNDERSTAND YOUR DATA, AND YOU UNDERSTAND HOW R SEES YOUR DATA, YOUR ANALYSIS WILL BE MUCH EASIER AND MORE EFFECTIVE**
 - Some questions you might have had before this session may include those on the slide:
 - How can I read data in R?
 - What are the basic data types in R?
 - How do I represent categorical information in R?
 - This session aims to help you answer those
-

Learning Objectives

- In this section, you'll be learning about the data types in R:
 - **WHAT DATA IS**
 - You'll be learning about data *structures*:
 - **WHAT DATA IS BUILT INTO - HOW IT IS ARRANGED**
 - You'll also learn how to find out what type/structure a particular piece of data has
 - Putting it together, you'll see how R's data types and structures relate to the types of data that you work with, yourself.
 - This will help you work much more fluently in R
 - In particular, we'll be working with **DATA FRAMES** - this is the workhorse data structure in R and, when you use R, you will see a *lot* of them.
-

Data Types and Structures in R

- R is **MOSTLY USED FOR DATA ANALYSIS**
- R is set up with key, core data types designed to help you work with your own data

- A lot of the time, R focuses on tabular data.
- **R is very powerful when dealing with tabular data**
- **INTERACTIVE DEMO**
- **SWITCH TO THE CONSOLE**
- Let's start by making a toy dataset.
- We'll eventually save this in your data/ directory, with the name feline-data.csv
- Use the console in RStudio
- We'll create a new variable called cats

- This holds a data.frame

```
cats <- data.frame(coat = c("calico", "black", "tabby"),
                  weight = c(2.1, 5.0, 3.2),
                  likes_catnip = c(1, 0, 1))
```

- We can now save the content of the variable cats as a CSV (comma-separated variable) file.
 - We use the write.csv() function
 - It is useful to call argument names explicitly so that the code is more readable

```
write.csv(cats, file = "data/feline-data.csv", row.names = FALSE)
```

- We're writing the cats data out to file
 - file specifies the filename we're writing to
 - row.names tells R not to give the rows in the table a name
- **USE THE FILES TAB/VIEW FILES TO VIEW THE CONTENTS OF THE NEW FILE**
- This is a plain text file, and you could have created it using a text editor like Nano or Notepad.
 - The first row is the header row
 - Each individual cat gets its own row
 - Each column is a different kind of data
 - Each column is separated by a comma
- We can load the data from a CSV file in a similar way to how we write it
 - We can use the read.csv() function
 - Inspect the dataset by using the variable name

```
cats <- read.csv(file = "data/feline-data.csv")
cats
```

	coat	weight	likes_catnip
1	calico	2.1	1
2	black	5.0	0
3	tabby	3.2	1

- **THINK ABOUT THE DATA TYPES** Are they all the same?
 - **NO** coat is text; weight is some real value (in kg or pounds, maybe), and likes_string looks like it should be TRUE/FALSE but is represented as 1s and 0s
 - **DOES IT MAKE SENSE TO WORK WITH EACH OF THESE ELEMENTS OF DATA AS IF THEY'RE THE SAME THING?** (No)
- Let's explore our dataset
- **EXTRACT A COLUMN FROM A TABLE**
 - Use \$ notation in the console
 - **NOTE THE AUTOCOMPLETION**

```
> cats$weight
[1] 2.1 5.0 3.2
> cats$coat
[1] "calico" "black" "tabby"
```

- **WHAT DID R RETURN?**
 - A *vector* (1D ordered collection) of numbers or character strings
- **WE CAN OPERATE ON THESE VECTORS**
 - *Vectors* are an important concept, and R is largely built so that operations on vectors are central to data analysis.
- Suppose that we discover the scales used to weigh the cats is miscalibrated
 - Every cat weight is underestimated by 2kg
 - We can correct all the weights at once

```
> cats$weight + 2
[1] 4.1 7.0 5.2
```

- Suppose we want to make a string out of each of the coat types:

```
> paste("My cat is", cats$coat)
[1] "My cat is calico" "My cat is black" "My cat is tabby"
```

- **What if we try to combine these columns?**

```
> cats$weight + cats$coat
Error in cats$weight + cats$coat :
non-numeric argument to binary operator
```

- **WE HIT AN ERROR**
 - You probably already realised that wasn't going to work, because adding "calico" to "2.1" is nonsense.
 - **You already have intuition about this**

- **THESE DATA TYPES ARE NOT COMPATIBLE** for addition
 - R's data types reflect the ways in which data is expected to interact
 - **UNDERSTANDING HOW YOUR OWN DATA MAP TO R's DATA TYPES IS KEY**
 - It's very important to understand how R sees your data (**you want R to see your data the same way you do**)
 - Many problems in R come down to incompatibilities between data and data types.
-

What Data Types Do You Expect?

- **ASK THE STUDENTS**
 - What data types would you expect to see?
 - What data types do you think you would **WANT OR NEED**, from your own experience?
 - **SPEND A COUPLE OF MINUTES ON THIS**
 - The difference between a *data type* and a *data structure*
-

Data Types in R

- R's data *types* are *atomic*: they are **FUNDAMENTAL AND EVERYTHING ELSE IS BUILT UP FROM THEM**, the same way matter is built up from atoms
 - In particular, all the data *structures* are built up from data *types*
- There are only **FIVE DATA TYPES** in R (though one is split into two...) - three that matter here and two we won't deal with
 - **logical**: Boolean, True/False (also 1/0)
 - **numeric**: anything that's a number on the number line; two types of number are supported: integer and double (real)
 - **character**: text data - readable symbols
- The other two are less relevant to you and we'll not deal with them much today
 - **complex**: complex numbers, defined on the 2D plane
 - **raw**: binary data
- **LET'S LEARN A BIT MORE ABOUT THEM IN THE DEMO**
- We can ask what type of data something is with the `typeof()` function

```
> typeof(cats$weight)
[1] "double"
> typeof(TRUE)
[1] "logical"
> typeof(3.14)
```

```
[1] "double"
> typeof(1)
[1] "double"
```

- This might seem unintuitive: 1 is an integer
 - But R treats all numbers as doubles (i.e. “Real” numbers) by default
 - To force R to see a value as an integer, we need to add an L suffix

```
> typeof(1L)
[1] "integer"
> typeof(cats$coat)
[1] "character"
> typeof("Irn Bru")
[1] "character"
```

- No matter how complicated your analysis gets, *all* data in R is interpreted as one of these basic data types.
 - R - like most computer languages - is strict about data types, and this has important consequences.

- **SUPPOSE OUR FRIEND WANTS TO ADD DETAILS OF THEIR OWN CAT TO OUR DATASET**

- We can create a new dataframe to hold this data

```
> additional_cat <- data.frame(coat = "tabby", weight = "2.3 or 2.4",
likes_catnip = 1)
> additional_cat
  coat      weight likes_catnip
1 tabby 2.3 or 2.4           1
```

- **We can “stack” dataframes using the function rbind**

```
> cats2 <- rbind(cats, additional_cat)
> cats2
  coat      weight likes_catnip
1 calico      2.1           1
2 black        5           0
3 tabby      3.2           1
4 tabby 2.3 or 2.4           1
```

- Initially, everything looks OK - but we’ve silently caused some problems.
 - Let’s look at the data type of cats\$weight in the two dataframes

```
> typeof(cats$weight)
[1] "double"
> typeof(cats2$weight)
[1] "character"
```

- Our friend gave us a different data type in the weight column, and this means we can no longer do the same weight adjustment (adding on 2kg) that we did before:

```
> cats2$weight + 2
Error in cats2$weight + 2 : non-numeric argument to binary operator
```

A Quick Note About Dataframes

- Any column in a dataframe can contain only one datatype.
- Initially, the datatype of `cats$weight` was double
- When we added the new cat, the data in that column was character - i.e. a string
- Because we can always represent a number as a string, but we can't always represent a string as a number, R chose to *coerce* the datatype from double to character for that column.
- **INTERACTIVE DEMO**
- We can look at the structure of a dataframe using the `str()` function

```
> str(cats)
'data.frame': 3 obs. of 3 variables:
 $ coat      : chr  "calico" "black" "tabby"
 $ weight    : num  2.1 5 3.2
 $ likes_catnip: int  1 0 1
> str(cats2)
'data.frame': 4 obs. of 3 variables:
 $ coat      : chr  "calico" "black" "tabby" "tabby"
 $ weight    : chr  "2.1" "5" "3.2" "2.3 or 2.4"
 $ likes_catnip: num  1 0 1 1
```

- Here the `str()` function tells us that `cats` and `cats2` are both `data.frames` - a very common kind of data structure in R
 - All data frames are composed of rows and columns.
 - Every column has the same number of rows
 - Every column has one and only one data type
 - Each column can be a different data type
- To understand `data.frames`, a bit better, let's meet a different data structure called a *vector*

Vectors

- These are the **MOST COMMON DATA STRUCTURE**
 - Vectors are an ordered collection of data values
 - Vectors can contain **ONLY A SINGLE DATA TYPE** (*atomic vectors*)
- **INTERACTIVE DEMO**
- Let's define an **ATOMIC VECTOR OF NUMBERS**

- To create a vector **USE THE c() FUNCTION** (c() is combine; use ?c)

```
> x <- c(10, 12, 45, 33)
> x
[1] 10 12 45 33
```

- Let's check the data type, and what kind of structure we have

```
> typeof(x)
[1] "double"
> length(x)
[1] 4
> str(x)
num [1:4] 10 12 45 33
```

- This is a little cryptic as output
 - We can see from typeof() that the vector contains the double type
 - The output of str() tells us:
 - it's a numeric (num) vector - which includes double and integer datatypes
 - there are four elements [1:4] in the vector
 - the first few datapoint in the vector
- str() tells us that the cats\$coat column is a vector, too:

```
> str(cats$coat)
chr [1:3] "calico" "black" "tabby"
```

Coercion (1)

- **INTERACTIVE DEMO**

- Given what we've done so far, what do you think the following will produce when we use typeof()?

- **PAUSE FOR STUDENT SUGGESTIONS**

```
> quiz_vector <- c(2,6,'3')
> typeof(quiz_vector)
[1] "character"
```

- Here, R has enforced that the type of the vector is character (string), because we can always represent numbers as strings, but we can't always represent strings as numbers
- **PAUSE: NEXT CALLOUTS**
- This is called *type coercion* and can cause surprises in your code

- It is a key reason why you need to be aware of the basic data types and how R interprets them.

- **INTERACTIVE DEMO**

- Consider these two vectors

```
> coercion_vector <- c('a', TRUE)
> coercion_vector
[1] "a"      "TRUE"
> another_coercion_vector <- c(0, TRUE)
> another_coercion_vector
[1] 0 1
```

- What do you expect their datatypes to be?

```
> typeof(coercion_vector)
[1] "character"
> typeof(another_coercion_vector)
[1] "double"
```

Coercion (2)

- *Coercion* is what happens when you **CONVERT ONE DATA TYPE INTO ANOTHER**
- If R thinks it needs to, it will **COERCE DATA IMPLICITLY** without telling you
- There is a set order for coercion
 - logical can be coerced to integer, but integer cannot be coerced to logical
 - That's because integer can describe all logical values, but not *vice versa*
 - Everything can be represented as a character, so that's the fallback position for R
- **IF THERE'S A FORMATTING PROBLEM IN YOUR DATA, R MIGHT CONVERT THE TYPE TO COPE**
 - R will choose the simplest data type that can represent all items in the vector
- **INTERACTIVE DEMO IN CONSOLE** More useful things to do with vectors
- You can (usually) **COERCE VECTORS MANUALLY** with `as.<type>()`

```
> x
[1] 10 12 45 33
> as.character(x)
[1] "10" "12" "45" "33"
> as.complex(x)
[1] 10+0i 12+0i 45+0i 33+0i
> as.logical(x)
[1] TRUE TRUE TRUE TRUE
```

- Surprising things can happen when R forces one data type into another.

- If your data doesn't look how you expect it to, *type coercion* may be to blame
 - Check your data formatting (is there an accidental string/character?)
- Let's look at our cats dataframe again

```
> cats
  coat weight likes_catnip
1 calico   2.1           1
2 black    5.0           0
3 tabby    3.2           1
> typeof(cats$likes_catnip)
[1] "integer"
```

- The type of the likes_catnip column is recorded as an integer, but we want logical TRUE/FALSE values.
 - We can use the as.logical() function to change this

```
> cats$likes_catnip <- as.logical(cats$likes_catnip)
> cats
  coat weight likes_catnip
1 calico   2.1         TRUE
2 black    5.0        FALSE
3 tabby    3.2         TRUE
> str(cats)
'data.frame': 3 obs. of  3 variables:
 $ coat      : chr  "calico" "black" "tabby"
 $ weight    : num  2.1 5 3.2
 $ likes_catnip: logi  TRUE FALSE TRUE
```

Lists

- lists are data structures like *vectors*, **EXCEPT THEY CAN HOLD ANY DATA TYPE**
 - They are not constrained to atomic types
 - They do not coerce their contents' datatypes
- Let's create a new list using the list() function

```
> l <- list(1, 'a', TRUE, seq(2, 5))
> length(l)
[1] 4
> l
[[1]]
[1] 1

[[2]]
[1] "a"

[[3]]
[1] TRUE
```

```
[[4]]  
[1] 2 3 4 5
```

- Individual elements in the list are identified with **DOUBLE SQUARE BRACKETS**
 - There are four elements
 - `[[1]]` is the number 1
 - `[[2]]` is the character "a"
 - `[[3]]` is the logical value TRUE
 - `[[4]]` is the vector of values from 2 to 5
- We can extract a single element from the list with the double square bracket notation

```
> l[[1]]  
[1] 1  
> l[[4]]  
[1] 2 3 4 5
```

- Using the `str()` function we can see the datatypes of all the elements in the list

```
> str(l)  
List of 4  
 $ : num 1  
 $ : chr "a"  
 $ : logi TRUE  
 $ : int [1:4] 2 3 4 5
```

- **The elements of a list can also have names**

- We specify the name when we create the list

```
> l_named <- list(a = "SWC", b = 1:4)  
> l_named  
$a  
[1] "SWC"  
  
$b  
[1] 1 2 3 4
```

- We can use the name of each element to retrieve it, with the `$` notation:

```
> l_named$a  
[1] "SWC"  
> l_named$b  
[1] 1 2 3 4
```

- But it's still a list like any other, and we can use the double square bracket notation.

```
> l_named[[1]]  
[1] "SWC"  
> l_named[[2]]  
[1] 1 2 3 4  
> str(l_named)  
List of 2
```

```
$ a: chr "SWC"
$ b: int [1:4] 1 2 3 4
```

Let's Look At A Data Frame

- We didn't go into detail about the cats dataframe we created at the start of this episode.
 - But let's look at it more closely

```
> cats
  coat weight likes_catnip
1 calico   2.1         TRUE
2 black    5.0        FALSE
3 tabby    3.2         TRUE
> typeof(cats)
[1] "list"
> cats[[2]]
[1] 2.1 5.0 3.2
> typeof(cats$weight)
[1] "double"
```

- So cats is a list, and each element in the list is a vector
- **PAUSE: NEXT CALLOUT**
- But cats is a **special kind of list** - a `data.frame` - where all the vectors have the same length.
 - We can see that this is a special kind of list by using the `class()` function

```
> class(cats)
[1] "data.frame"
> class(1)
[1] "list"
```

- The class `data.frame` represents a standard way of organising data:
 - Each *row* is an observation
 - Each *column* is a variable
 - The data frame represents a series of observations
-

Load Episode Data

- **DEMO IN SCRIPT**
- **DOWNLOAD DATA**
 - Use the link from the Etherpad document
 - Place the file in `data/`

- **CREATE A NEW SCRIPT**

- Call it gapminder
- Add the code
- Use read.table
- The data is in CSV format
- We need to provide a data source (**here, a file**), the separator character, and whether there's a header row

Load gapminder data from a local file

```
gapminder <- read.table("data/gapminder_data.csv", sep=",", header=TRUE)
```

- **RUN THE SCRIPT** (use Source)
- **CHECK THE DATA IN THE Environment TAB**
 - Click on gapminder in Environment tab.
 - **NOTE COLUMNS**

Investigating gapminder

- Now we've loaded our data, let's take a look at it
- **DEMO IN CONSOLE**
 - 1704 rows, 6 columns
 - Investigate types of columns
 - **POINT OUT THAT THE TYPE OF A COLUMN IS INTEGER IF IT'S A FACTOR**
 - **LENGTH OF A DATAFRAME IS THE NUMBER OF COLUMNS**
- It's always useful to get an initial understanding of your data with the str() function.

```
> str(gapminder)
'data.frame': 1704 obs. of  6 variables:
 $ country  : Factor w/ 142 levels "Afghanistan",...: 1 1 1 1 1 1 1 1 1 1 ...
 $ year     : int   1952 1957 1962 1967 1972 1977 1982 1987 1992 1997 ...
 $ pop      : num   8425333 9240934 10267083 11537966 13079460 ...
 $ continent: Factor w/  5 levels "Africa","Americas",...: 3 3 3 3 3 3 3 3 3 3
...
 $ lifeExp  : num   28.8 30.3 32 34 36.1 ...
 $ gdpPercap: num   779 821 853 836 740 ...
```

- The summary() function can also be useful
 - With dataframes, this gives a numeric, tabular, or descriptive summary of each column.

```
> summary(gapminder)
      country      year      pop      continent
lifeExp
Afghanistan: 12   Min.   :1952   Min.   :6.001e+04   Africa   :624   Min.
:23.60
```

```

Albania      : 12 1st Qu.:1966 1st Qu.:2.794e+06 Americas:300 1st
Qu.:48.20
Algeria      : 12 Median :1980 Median :7.024e+06 Asia :396 Median
:60.71
Angola       : 12 Mean :1980 Mean :2.960e+07 Europe :360 Mean
:59.47
Argentina    : 12 3rd Qu.:1993 3rd Qu.:1.959e+07 Oceania : 24 3rd
Qu.:70.85
Australia    : 12 Max. :2007 Max. :1.319e+09 Max.
:82.60
(Other)      :1632
gdpPercap
Min. : 241.2
1st Qu.: 1202.1
Median : 3531.8
Mean : 7215.3
3rd Qu.: 9325.5
Max. :113523.1

```

- We can also inspect individual columns of the data.

```

> typeof(gapminder$year)
[1] "integer"
> typeof(gapminder$country)
[1] "integer"

```

- This looks a little weird: why should the country be an integer?
 - Let's investigate further with str()

```

> str(gapminder$country)
Factor w/ 142 levels "Afghanistan",...: 1 1 1 1 1 1 1 1 1 1 ...

```

- The country column has been read in as a factor
 - A factor is a data structure that represents *categorical* data.
 - We can see this in the str() output we got, as well: both country and continent were read in as factors.
 - This is fine, and is what we actually want.

- We can ask about the dimensions of the dataframe with dim()

```

> dim(gapminder)
[1] 1704 6

```

- So there are 1704 rows and 6 columns
- What does length() produce:

```

> length(gapminder)
[1] 6

```

- Remember that a dataframe is a list of vector columns, so its length is the number of elements in the list.
 - There are six columns, so the length of the dataframe is 6.

- We can ask specifically for the number of rows and columns, and for other information.
- We can also use `head()` to summarise the dataframe.

```
> nrow(gapminder)
[1] 1704
> ncol(gapminder)
[1] 6
> colnames(gapminder)
[1] "country" "year" "pop" "continent" "lifeExp" "gdpPercap"
> head(gapminder)
```

	country	year	pop	continent	lifeExp	gdpPercap
1	Afghanistan	1952	8425333	Asia	28.801	779.4453
2	Afghanistan	1957	9240934	Asia	30.332	820.8530
3	Afghanistan	1962	10267083	Asia	31.997	853.1007
4	Afghanistan	1967	11537966	Asia	34.020	836.1971
5	Afghanistan	1972	13079460	Asia	36.088	739.9811

Data Frame Manipulation With `dplyr`

Learning Objectives (Data Frames)

- You're going to **learn to manipulate data.frames with the six verbs of dplyr**
- `select()`
- `filter()`
- `group_by()`
- `summarize()`
- `mutate()`
- `%>%` (pipe)

What and Why is `dplyr`?

- `dplyr` is a package in the **TIDYVERSE**; it exists to enable **rapid analysis of data by groups**
 - For example, if we wanted numerical (rather than graphical) analysis of the `gapminder` data by continent, we'd use `dplyr`
 - It enables group-level analyses without using repetitive code
- **AVOIDING REPETITION IMPROVES YOUR CODE**

- More **robust**
 - More **readable**
 - More **reproducible**
-

Split-Apply-Combine

- The **general principle** dplyr supports is called **SPLIT-APPLY-COMBINE**
 - We have a **dataset with several groups in a variable** (column x)
 - For example, each patient in our messy data might be a “group”
 - We **want to perform the same operation on each group, independently** - take a mean of y for each group, for example
 - So we **SPLIT** the data into groups, on x
 - Then we **APPLY** the operation (take the mean for each group)
 - Then we **COMBINE** the results into a new table
-

select() - Interactive Demo**

- **DEMO IN CONSOLE**

- Import dplyr

```
> library(dplyr)
```

- The select() *verb* **SELECTS COLUMNS**

- **DEMO IN CONSOLE**

- If we wanted to select only year, country and GDP data from gapminder

- Specify: **data, then columns**

```
> head(select(gapminder, year, country, gdpPercap))
```

```
year    country gdpPercap
1 1952 Afghanistan  779.4453
2 1957 Afghanistan  820.8530
3 1962 Afghanistan  853.1007
4 1967 Afghanistan  836.1971
5 1972 Afghanistan  739.9811
6 1977 Afghanistan  786.1134
```

- Here, we **applied a function**, but we can also **‘PIPE’ DATA FROM ONE VERB TO ANOTHER**

- These work **like pipes in the shell**

- **SPECIAL PIPE SYMBOL: %>%**

- Specify **only columns**

```
> gapminder %>% select(year, country, gdpPercap) %>% head()
year    country gdpPercap
```

```
1 1952 Afghanistan 779.4453
2 1957 Afghanistan 820.8530
3 1962 Afghanistan 853.1007
4 1967 Afghanistan 836.1971
5 1972 Afghanistan 739.9811
6 1977 Afghanistan 786.1134
```

`filter()`

- `filter()` selects rows on the basis of some condition, or combination of conditions
 - We can **use it as a function, with pipes**

- **DEMO IN CONSOLE**

```
> head(filter(gapminder, continent=="Europe"))
  country year      pop continent lifeExp gdpPercap
1 Albania 1952 1282697    Europe   55.23  1601.056
2 Albania 1957 1476505    Europe   59.28  1942.284
3 Albania 1962 1728137    Europe   64.82  2312.889
4 Albania 1967 1984060    Europe   66.22  2760.197
5 Albania 1972 2263554    Europe   67.69  3313.422
6 Albania 1977 2509048    Europe   68.93  3533.004
```

- **DEMO IN SCRIPT** (gapminder.R)

- One **advantage of pipes** is that they make chaining *verbs* together **MORE READABLE**
- **END THE LINES WITH THE PIPE SYMBOL** so R knows that there's a continuation
- Run the lines and **check the output** in Environment

```
# Select gdpPercap by country and year, only for Europe
eurodata <- gapminder %>%
  filter(continent == "Europe") %>%
  select(year, country, gdpPercap)
```

Challenge

```
# Select life expectancy by country and year, only for Africa
> afrodata <- gapminder %>%
  filter(continent == "Africa") %>%
  select(year, country, lifeExp)
> nrow(afrodata)
[1] 624
```



Red sticky for a question or issue



Green sticky if complete



group_by()

- The group_by() verb **SPLITS** data.frames INTO GROUPS ON A **VARIABLE/COLUMN PROPERTY**
- **DEMO IN CONSOLE**
 - It returns a **tibble** - a table with extra metadata describing the groups in the table

```
> group_by(gapminder, continent)
# A tibble: 1,704 x 6
# Groups:   continent [5]
  country    year    pop continent lifeExp gdpPercap
  <fctr> <int>   <dbl>   <fctr>   <dbl>   <dbl>
1 Afghanistan 1952  8425333   Asia  28.801  779.4453
2 Afghanistan 1957  9240934   Asia  30.332  820.8530
3 Afghanistan 1962 10267083   Asia  31.997  853.1007
4 Afghanistan 1967 11537966   Asia  34.020  836.1971
5 Afghanistan 1972 13079460   Asia  36.088  739.9811
6 Afghanistan 1977 14880372   Asia  38.438  786.1134
7 Afghanistan 1982 12881816   Asia  39.854  978.0114
8 Afghanistan 1987 13867957   Asia  40.822  852.3959
9 Afghanistan 1992 16317921   Asia  41.674  649.3414
10 Afghanistan 1997 22227415   Asia  41.763  635.3414
# ... with 1,694 more rows
```

summarize()

- The **combination of group_by() and summarize() is very powerful**
 - We can **CREATE NEW VARIABLES** using functions that repeat for each group
- Here, we've split the original table into three groups, and now **CREATE A NEW VARIABLE mean_b THAT IS FILLED BY CALCULATING THE MEAN OF b**
- **DEMO IN SCRIPT**
 - We use the same principle to **calculate mean GDP per continent**

```
> # Produce table of mean GDP by continent
> gapminder %>%
+   group_by(continent) %>%
+   summarize(meangdpPercap=mean(gdpPercap))
# A tibble: 5 x 2
  continent meangdpPercap
  <fctr>      <dbl>
1 Africa      2193.755
2 Americas    7136.110
3 Asia        7902.150
```

```
4 Europe 14469.476
5 Oceania 18621.609
```

Challenge 13

- **IN THE SCRIPT**

```
# Find average life expectancy by nation
avg_lifexp_country <- gapminder %>%
  group_by(country) %>%
  summarize(meanlifeExp=mean(lifeExp))
```

- **IN THE CONSOLE**

```
> avg_lifexp_country %>% filter(meanlifeExp == min(meanlifeExp))
# A tibble: 1 × 2
  country      meanlifeExp
  <chr>         <dbl>
1 Sierra Leone 36.8
> avg_lifexp_country %>% filter(meanlifeExp == max(meanlifeExp))
# A tibble: 1 × 2
  country meanlifeExp
  <chr>         <dbl>
1 Iceland 76.5
```



Red sticky for a question or issue



Green sticky if complete



count() and n()

- Two other useful functions are related to summarize()
 - **count()** reports a new table of counts by group
 - **n()** is used to represent the count of rows, when calculating new values in summarize()
- **DEMO IN CONSOLE**
- **NOTE:** standard error is (std dev)/sqrt(n)

```
> gapminder %>% filter(year == 2002) %>% count(continent, sort = TRUE)
# A tibble: 5 × 2
  continent      n
  <fctr> <int>
1 Africa 52
2 Asia 33
3 Europe 30
```

```

4 Americas      25
5 Oceania       2
> gapminder %>% group_by(continent) %>% summarize(se_lifeExp =
sd(lifeExp)/sqrt(n()))
# A tibble: 5 x 2
  continent se_lifeExp
  <fctr>    <dbl>
1 Africa   0.3663016
2 Americas 0.5395389
3 Asia     0.5962151
4 Europe   0.2863536
5 Oceania  0.7747759

```

`mutate()`

- **`mutate()` CALCULATES NEW VARIABLES (COLUMNS) ON THE BASIS OF EXISTING COLUMNS**
- **DEMO IN SCRIPT**
 - Say we want to calculate the **total GDP of each nation, each year, in \$bn**
 - We'd multiply the GDP per capita by the total population, and divide by 1bn
- **INSPECT THE OUTPUT**
 - We have a new data table, which is the gapminder data, plus an extra column

Calculate GDP in \$billion

```

gdp_bill <- gapminder %>%
  mutate(gdp_billion = gdpPercap * pop / 10^9)

```

- **WE CAN CHAIN ALL THESE OPERATIONS TOGETHER WITH PIPES**
- **We can calculate several summaries in a single `summarize()` command**
- We can use the output of `mutate()` in the `summarize()` command
- **DEMO IN SCRIPT**
 - We're going to calculate the **total (and standard deviation) of GDP per continent, per year**
 - Calculate total GDP first
 - Group by continent and year
 - Summarise mean and sd of GDP per capita, and total GDP
- **INSPECT THE OUTPUT**

Calculate total/sd of GDP by continent and year

```

gdp_bycontinents_byyear <- gapminder %>%
  mutate(gdp_billion=gdpPercap*pop/10^9) %>%
  group_by(continent,year) %>%
  summarize(mean_gdpPercap=mean(gdpPercap),
            sd_gdpPercap=sd(gdpPercap),
            mean_gdp_billion=mean(gdp_billion),
            sd_gdp_billion=sd(gdp_billion))

```

```
ifelse()
```

- **ifelse() IS A FILTER THAT CAN BE USED WITH MUTATE TO CALCULATES NEW VARIABLES (COLUMNS) ON THE BASIS OF EXISTING COLUMNS ONLY IF SOME CONDITION IS MET**
- **DEMO IN SCRIPT**
 - Say we want to calculate the total GDP of each nation, each year, in \$bn **BUT ONLY FOR COUNTRIES OVER 10mn PEOPLE**

```
gdp_billion_large_countries <- gapminder %>%  
  mutate(gdp_billion_large = ifelse(pop > 10e6,  
                                     gdpPercap * pop / 10^9,  
                                     NA))
```

- **INSPECT THE OUTPUT**
 - We have a new data table, which is the gapminder data, plus an extra column
- Similarly, we can scale GDP only in those cases where life expectancy for a nation is over 40

```
gdp_future_bycontinents_byyear_high_lifeExp <- gapminder %>%  
  mutate(gdp_futureExpectation = ifelse(lifeExp > 40, gdpPercap * 1.5,  
                                         gdpPercap)) %>%  
  group_by(continent, year) %>%  
  summarize(mean_gdpPercap = mean(gdpPercap),  
             mean_gdpPercap_expected = mean(gdp_futureExpectation))
```

Tidy Data

Why Tidy Data?

- Data cleaning/processing is **not just a first step** - it must be repeated many time over the course of an analysis
 - new data, new ideas, etc. turn up as you're working
 - About **80% of the effort of data analysis** is cleaning and preparing data for analysis
 - The principles of **tidy data** provide a standard way to organise data values within a dataset
 - "Tidy datasets are all alike, but every messy dataset is messy in its own way"
-

An Untidy Dataset (1)

- Here's a dataset like you might receive it from a colleague
 - It's a **RECTANGULAR TABLE**
 - Made up of **ROWS** and **COLUMNS**, just like the data you've been working with
 - Each row describes a treatment
 - Each column gives the results for a different individual, for each treatment
 - Thinking about how our dataframes are structured, there should be one row per observation.
 - But here the observations are the people
 - So maybe we could transpose the rows and columns so that the people are the rows?
-

An Untidy Dataset (2)

- So we've transposed the rows and columns of the table
 - The **data is the same**
 - The **layout is different**
 - Now we have one row per observation, and we have one column per variable (treatments A and B), and that's what we want, isn't it?
 - **PAUSE: NEW CALLOUT**
 - But the data doesn't have to be structured this way.
 - In fact, this isn't a very good way to structure data for many analyses.
 - To understand what this means, **AND WHY IT IS UNTIDY DATA**, we need to consider some data semantics.
-

Data Semantics

- We need to define three terms
- A dataset, like the one shown, is a collection of **VALUES**
- Each value belongs to a variable, and to an observation

- A **VARIABLE** is something that *can change or vary*
 - They may be values that measure the same underlying attribute, such as height, temperature, or some kind of output result
 - They may be experimental conditions under a researcher's control, such as a treatment, how long that treatment is applied, or other experimental settings
 - An **OBSERVATION** is a collection of **values** measured across all **variables** for the same individual or unit
 - The unit is often a person, a collective group like a religion or company, or a physical item like a reactor vessel
-

Challenge (2min)

- So for our first messy dataset, how would you describe the rows and columns of the table?
 - Are they observations, or variables, or neither?
 - **THE ROWS AND COLUMNS IN THE MESSY DATA ARE NEITHER OBSERVATIONS NOR VARIABLES**
-

A Tidy Dataset (1)

- The dataset contains 18 values:
 - six observations of three variables
 - The variables are:
 - **PERSON**: John, Mary and Jane
 - **TREATMENT**: A or B
 - **RESULT**: NA, 16, 3, 2, 11, 1
 - Each **OBSERVATION** includes values for all three variables
-

A Tidy Dataset (2)

- Here is the dataset represented in tidy form
 - You can see each variable has its own column
 - Each observation has one value per column
 - **Note: Missing data counts as a value**
-

Tidy Data (2)

- Tidy data is a **STANDARD** way of structuring a dataset, but it is not the only way, or always the best way
 - It does make it easy to extract the variables you need
 - It is also very well suited to R because it supports vectorisation (which you saw earlier)
- In Tidy Data:
 - Each variable forms a column
 - Each observation forms a row
- Most of the data you receive is unlikely to be Tidy, so you'll probably need to clean it
- And **MESSY DATA CAN BE USEFUL**
 - If your design is completely crossed: e.g. every individual tries every medicine
 - Messy Data (rows=patients, columns=medicines) is compact
 - Also useful if matrix operations are appropriate
- **INTERACTIVE DEMO**
 - Show that the gapminder data is in an intermediate format
 - Three ID variables (continent, country, year)
 - Three observation variables (pop, lifeExp, gdpPercap)
 - This intermediate form can be preferable; there is no advantage to having a single "observation" column, here
- **IN THE CONSOLE**
 - Each column in the gapminder dataset is a variable
 - Each row is a set of values: one per variable, comprising a single observation for a combination of country and year
- **WE CAN USE DPLYR DATA METHODS ON THIS DATASET**

```
> head(gapminder)
  country year      pop continent lifeExp gdpPercap
1 Afghanistan 1952  8425333      Asia  28.801  779.4453
2 Afghanistan 1957  9240934      Asia  30.332  820.8530
3 Afghanistan 1962 10267083      Asia  31.997  853.1007
4 Afghanistan 1967 11537966      Asia  34.020  836.1971
5 Afghanistan 1972 13079460      Asia  36.088  739.9811
6 Afghanistan 1977 14880372      Asia  38.438  786.1134
```

Long v Wide

- We often refer to datasets as being "long" or "wide"
- **LONG datasets have one row per observation, and one column per variable**

- **WIDE datasets might have multiple arrangements**
- Wide datasets tend to be easier for humans to read
 - Quite a lot of the time, we tend to record and share datasets in wide format.
- **BUT MANY R FUNCTIONS ARE DESIGNED TO WORK WITH LONG DATA**
 - Some plotting functions work better with wide data, though
 - The purely long format can require additional grouping operations (like `group_by()`), and it can be more convenient to have an intermediate form like the gapminder data.

gapminder Wide Dataset

- We've been using the nicely-formatted gapminder data to make our lives easier.
- **THE DATA YOU RECEIVE IN A RESEARCH CONTEXT WILL NOT USUALLY BE SO NICELY-FORMATTED**
- Let's load in the wide-format gapminder data and take a look at it.

```
> gap_wide <- read.csv("data/gapminder_wide.csv", stringsAsFactors = FALSE)
> str(gap_wide)
'data.frame': 142 obs. of  38 variables:
 $ continent      : chr  "Africa" "Africa" "Africa" "Africa" ...
 $ country        : chr  "Algeria" "Angola" "Benin" "Botswana" ...
 $ gdpPercap_1952 : num  2449 3521 1063 851 543 ...
 $ gdpPercap_1957 : num  3014 3828 960 918 617 ...
 $ gdpPercap_1962 : num  2551 4269 949 984 723 ...
 $ gdpPercap_1967 : num  3247 5523 1036 1215 795 ...
 $ gdpPercap_1972 : num  4183 5473 1086 2264 855 ...
 $ gdpPercap_1977 : num  4910 3009 1029 3215 743 ...
 $ gdpPercap_1982 : num  5745 2757 1278 4551 807 ...
 $ gdpPercap_1987 : num  5681 2430 1226 6206 912 ...
 $ gdpPercap_1992 : num  5023 2628 1191 7954 932 ...
 $ gdpPercap_1997 : num  4797 2277 1233 8647 946 ...
 $ gdpPercap_2002 : num  5288 2773 1373 11004 1038 ...
 $ gdpPercap_2007 : num  6223 4797 1441 12570 1217 ...
 $ lifeExp_1952   : num  43.1 30 38.2 47.6 32 ...
 $ lifeExp_1957   : num  45.7 32 40.4 49.6 34.9 ...
 $ lifeExp_1962   : num  48.3 34 42.6 51.5 37.8 ...
 $ lifeExp_1967   : num  51.4 36 44.9 53.3 40.7 ...
 $ lifeExp_1972   : num  54.5 37.9 47 56 43.6 ...
 $ lifeExp_1977   : num  58 39.5 49.2 59.3 46.1 ...
 $ lifeExp_1982   : num  61.4 39.9 50.9 61.5 48.1 ...
 $ lifeExp_1987   : num  65.8 39.9 52.3 63.6 49.6 ...
 $ lifeExp_1992   : num  67.7 40.6 53.9 62.7 50.3 ...
 $ lifeExp_1997   : num  69.2 41 54.8 52.6 50.3 ...
 $ lifeExp_2002   : num  71 41 54.4 46.6 50.6 ...
 $ lifeExp_2007   : num  72.3 42.7 56.7 50.7 52.3 ...
```



```

$ pop_1952      : num  9279525 4232095 1738315 442308 4469979 ...
$ pop_1957      : num  10270856 4561361 1925173 474639 4713416 ...
$ pop_1962      : num  11000948 4826015 2151895 512764 4919632 ...
$ pop_1967      : num  12760499 5247469 2427334 553541 5127935 ...
$ pop_1972      : num  14760787 5894858 2761407 619351 5433886 ...
$ pop_1977      : num  17152804 6162675 3168267 781472 5889574 ...
$ pop_1982      : num  20033753 7016384 3641603 970347 6634596 ...
$ pop_1987      : num  23254956 7874230 4243788 1151184 7586551 ...
$ pop_1992      : num  26298373 8735988 4981671 1342614 8878303 ...
$ pop_1997      : num  29072015 9875024 6066080 1536536 10352843 ...
$ pop_2002      : int   31287142 10866106 7026113 1630347 12251209 7021078
15929988 4048013 8835739 614382 ...
$ pop_2007      : int   33333216 12420476 8078314 1639131 14326203 8390505
17696293 4369038 10238807 710960 ...

```

- **ALSO VIEW IN RSTUDIO**

- This is very wide data and it looks like it would be a bit of a nightmare to do the kinds of analyses we have been doing, with the data in this format.

Pivot: Wide to Long

- We want to convert this wide format that's difficult to work with to our nice, longer, intermediate layout.
- **TO DO THIS WE USE THE `pivot_longer()` FUNCTION FROM `dplyr/tidyverse`**
 - This makes datasets longer by increasing the number of rows and decreasing the number of columns
 - You may be familiar with “pivot tables” from Excel, which are similar.
- It works like the image shows
 - We define the dataset we want to pivot
 - We say which columns we want to convert from wide to long format in the `cols` variable
 - The `pivot_longer()` function effectively splits the table by these columns
 - The column name then gets converted to its own, new column, but the values are kept
 - We specify the names of the new “name” and “value” columns with `names_to` and `values_to`.
- For the `gapminder` data, we want to pivot all columns that start with `pop`, `lifeExp` or `gdpPercent`
 - We'll put the column names into a column called `obstype_year`
 - We'll put the values into a column called `obs_values`

```

> gap_long <- gap_wide %>%
+   pivot_longer(
+     cols = c(starts_with('pop'), starts_with('lifeExp'),

```

```
starts_with('gdpPercap')),
+       names_to = "obstype_year", values_to = "obs_values"
+     )
> str(gap_long)
tibble [5,112 × 4] (S3: tbl_df/tbl/data.frame)
 $ continent   : chr [1:5112] "Africa" "Africa" "Africa" "Africa" ...
 $ country     : chr [1:5112] "Algeria" "Algeria" "Algeria" "Algeria" ...
 $ obstype_year: chr [1:5112] "pop_1952" "pop_1957" "pop_1962" "pop_1967" ...
 $ obs_values  : num [1:5112] 9279525 10270856 11000948 12760499 14760787 ...
```

- This gives us a long format table with four columns
- **BUT IT'S NOT QUITE WHAT WE WANT, BECAUSE WE HAVE TWO KINDS OF DATA COMBINED IN THE obstype_year COLUMN**
 - We want to split year into its own column.

`separate()`

- **THE separate() FUNCTION SPLITS THE CONTENTS OF A SINGLE COLUMN INTO MULTIPLE NEW COLUMNS**
 - We want to split the year information from every cell in the obstype_year column into its own new column.
- The separate() function needs to know:
 - dataframe it's working on
 - the column it's splitting
 - the names of the new columns it's making
 - the character it's splitting the cell content on
- We split the obstype_year column into two columns called obs_type and year, on the underscore separator "_".

```
> gap_long <- gap_long %>% separate(obstype_year, into = c('obs_type',
'year'), sep = "_")
```

- As the year at this point is still a character string, we change the datatype with `as.integer` and take a look at the new dataframe.

```
> gap_long$year <- as.integer(gap_long$year)
> str(gap_long)
tibble [5,112 × 5] (S3: tbl_df/tbl/data.frame)
 $ continent   : chr [1:5112] "Africa" "Africa" "Africa" "Africa" ...
 $ country     : chr [1:5112] "Algeria" "Algeria" "Algeria" "Algeria" ...
 $ obs_type    : chr [1:5112] "pop" "pop" "pop" "pop" ...
 $ year       : int [1:5112] 1952 1957 1962 1967 1972 1977 1982 1987 1992 1997
 ...
 $ obs_values  : num [1:5112] 9279525 10270856 11000948 12760499 14760787 ...
```

- **THIS IS STILL NOT QUITE WHERE WE WANT IT**
 - There are three types of observation combined in obs_type and obs_values and it would be convenient to split them into new columns

Pivot: Long to Wide

- We want to convert this long format into a slightly wider intermediate layout, for convenience.
- There is a companion function to `pivot_longer()` called `pivot_wider()` that we use for this.
 - It's essentially the reverse of `pivot_longer()`
- **THE `pivot_wider()` FUNCTION SPLITS THE TABLE UP BY VALUES IN THE COLUMN OF "NAMES"**
 - It then changes the name of the "VALUES" column to reflect the value in the "NAMES" column for each split.
 - Finally, it recombines the columns into a single table on the basis of the ID columns
- We want to do this with the `gapminder` data by using the `obs_type` column for the names, and the `obs_values` column for the values.

```
> gap_normal <- gap_long %>% pivot_wider(names_from = obs_type, values_from = obs_values)
> str(gap_normal)
tibble [1,704 × 6] (S3: tbl_df/tbl/data.frame)
 $ continent: chr [1:1704] "Africa" "Africa" "Africa" "Africa" ...
 $ country  : chr [1:1704] "Algeria" "Algeria" "Algeria" "Algeria" ...
 $ year     : int [1:1704] 1952 1957 1962 1967 1972 1977 1982 1987 1992 1997
 ...
 $ pop      : num [1:1704] 9279525 10270856 11000948 12760499 14760787 ...
 $ lifeExp  : num [1:1704] 43.1 45.7 48.3 51.4 54.5 ...
 $ gdpPercap: num [1:1704] 2449 3014 2551 3247 4183 ...
```

- And we can see that this has restored the intermediate form of the data that was so useful to us earlier.

Challenge (5min)

- Using `gap_long`, calculate the mean life expectancy, population, and `gdpPercap` for each continent. Hint: use the `group_by()` and `summarize()` functions we learned in the `dplyr` lesson

```
gap_long %>% group_by(continent, obs_type) %>%
  summarize(means=mean(obs_values))
```



Red sticky for a question or issue



Green sticky if complete

